

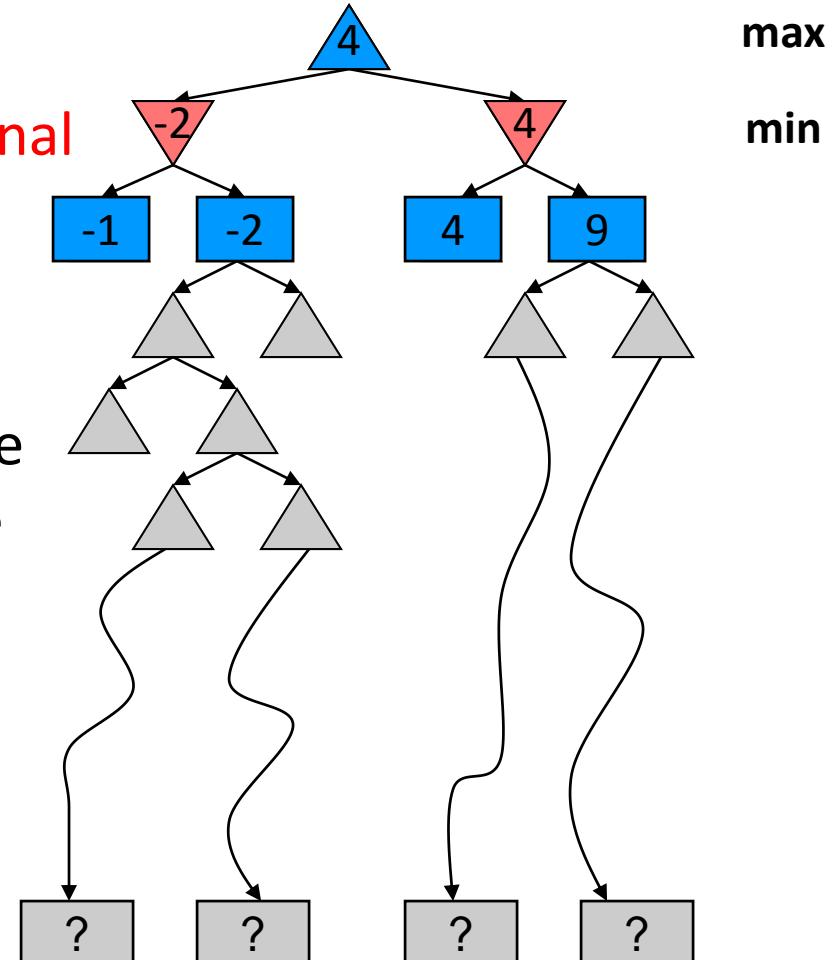
# Evaluation Functions

# Resource Limits

- The minimax algorithm generates the **entire game search space**, whereas the alpha–beta algorithm allows us to prune large parts of it.

- However, alpha–beta still has to **search all the way to terminal states for at least a portion of the search space**

- we turn to evaluation functions, functions that take in a state and output an estimate of the true minimax value of that node



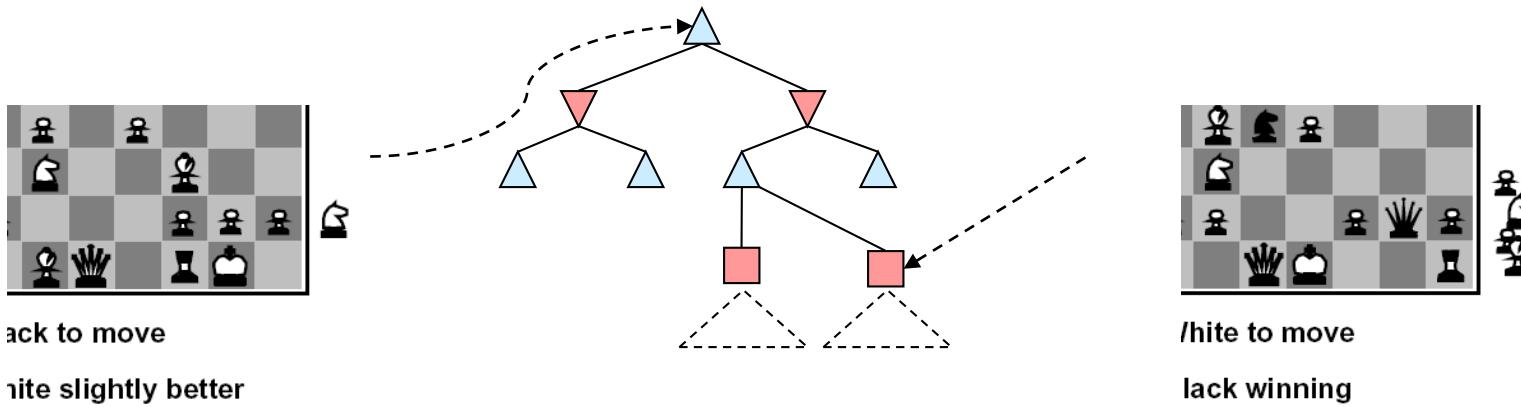
# Resource Limits

- Evaluation functions are widely employed in depth-limited minimax.
- where we treat non-terminal nodes located at our maximum solvable depth as terminal nodes.
- Assigning mock terminal utilities as determined by a carefully selected evaluation function.
- Replace terminal utilities with an evaluation function for non-terminal positions (Similar to heuristics in Informed search)

- Example:
  - Suppose we have 100 seconds, can explore 10K nodes / sec
  - So can check 1M nodes per move
  - $\beta$ - $\alpha$  reaches about depth 8 – decent chess program
  - what is depth 8 means in chess ?
- evaluation functions can only yield estimates of the values of non-terminal utilities .Guarantee of optimal play is gone
- More depth makes a BIG difference

# Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

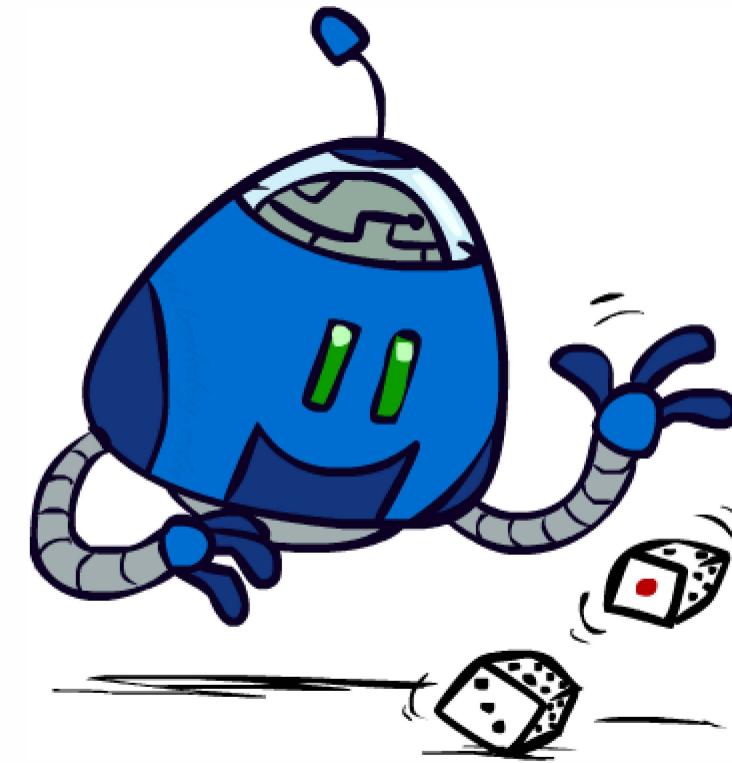
$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ , etc.

# Depth Matters

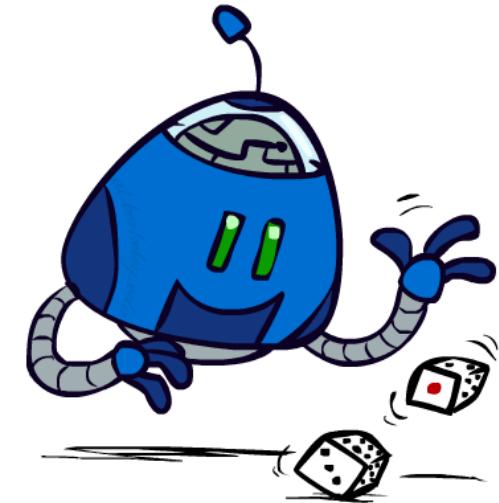
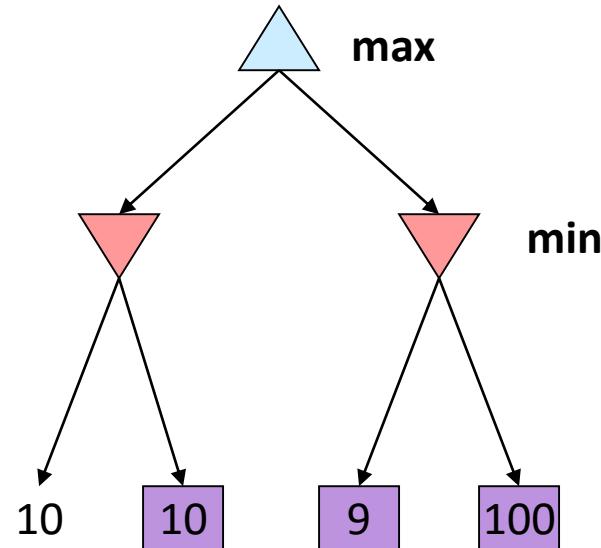
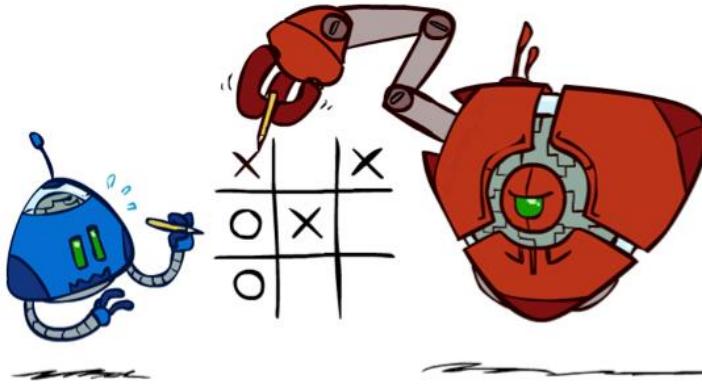
- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation

# Expectimax Search



# UNCERTAIN OUTCOMES

# Worst-Case vs. Average Case

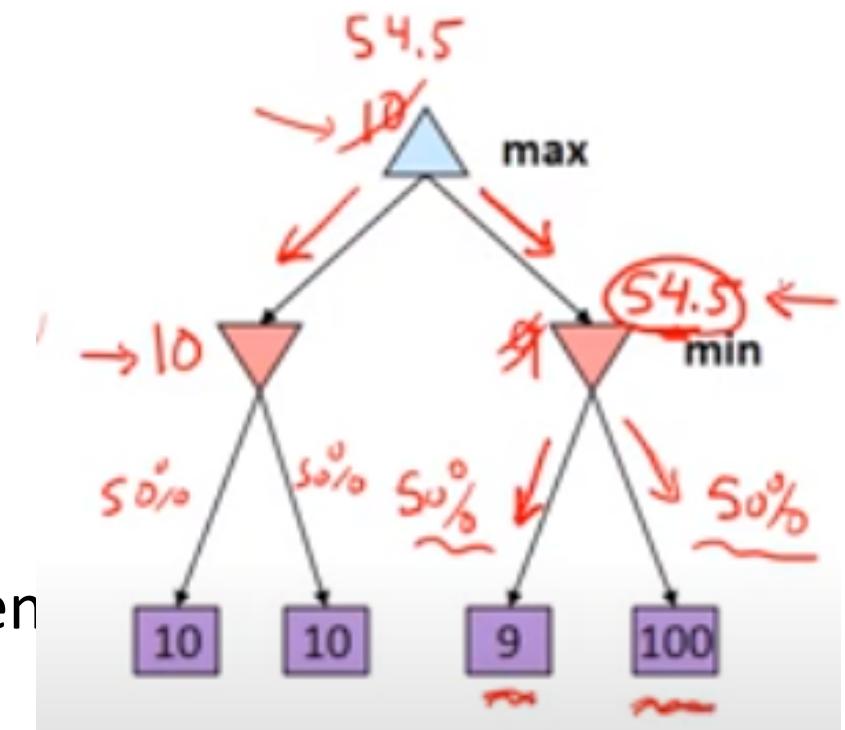
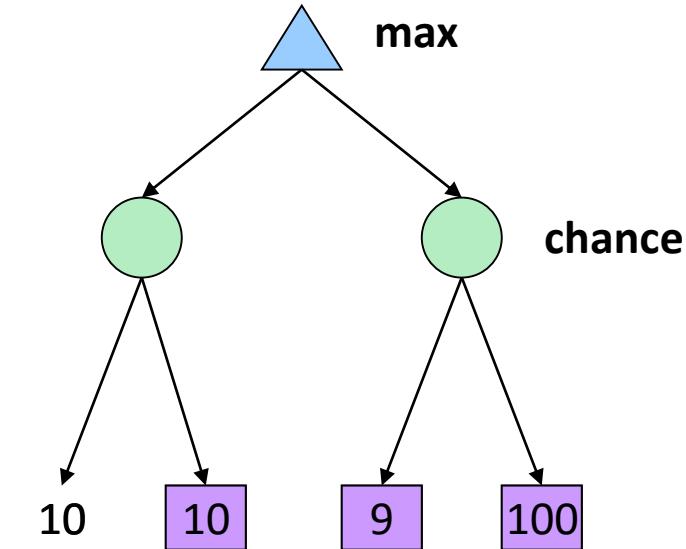


Idea: Uncertain outcomes controlled by chance, not an adversary!

In this example we will never get 100. we say never because we are assuming playing against optimal player. (**Worst case**)

# Expectimax Search

- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly
  - Actions can fail: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- Expectimax search: compute the average score under optimal play
  - Max nodes as in minimax search
  - Chance nodes are like min nodes but the outcome is uncertain
  - Calculate their expected utilities
    - i.e., take weighted average (expectation) of children

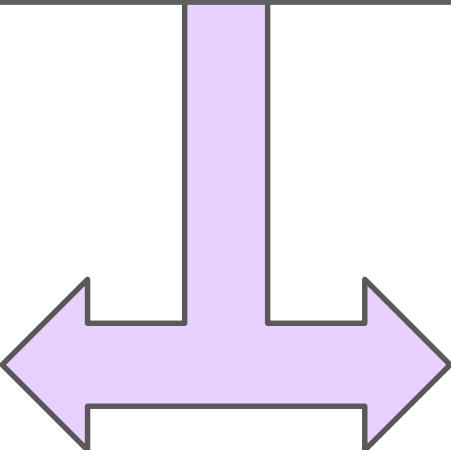


# Expectimax Pseudocode

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is EXP: return exp-value(state)
```

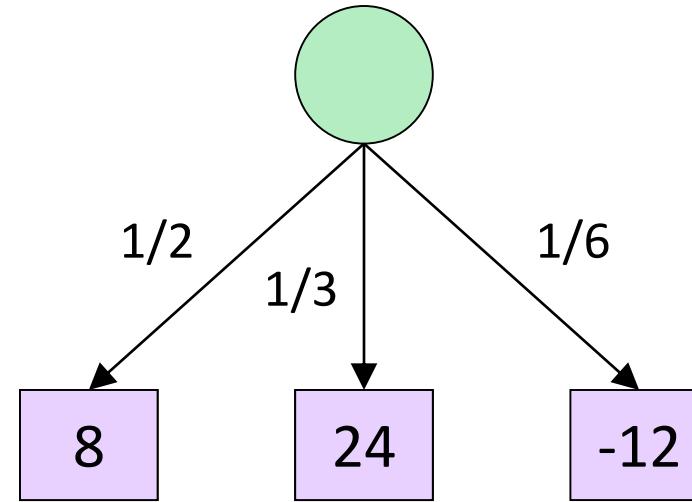
```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p * value(successor)
    return v
```



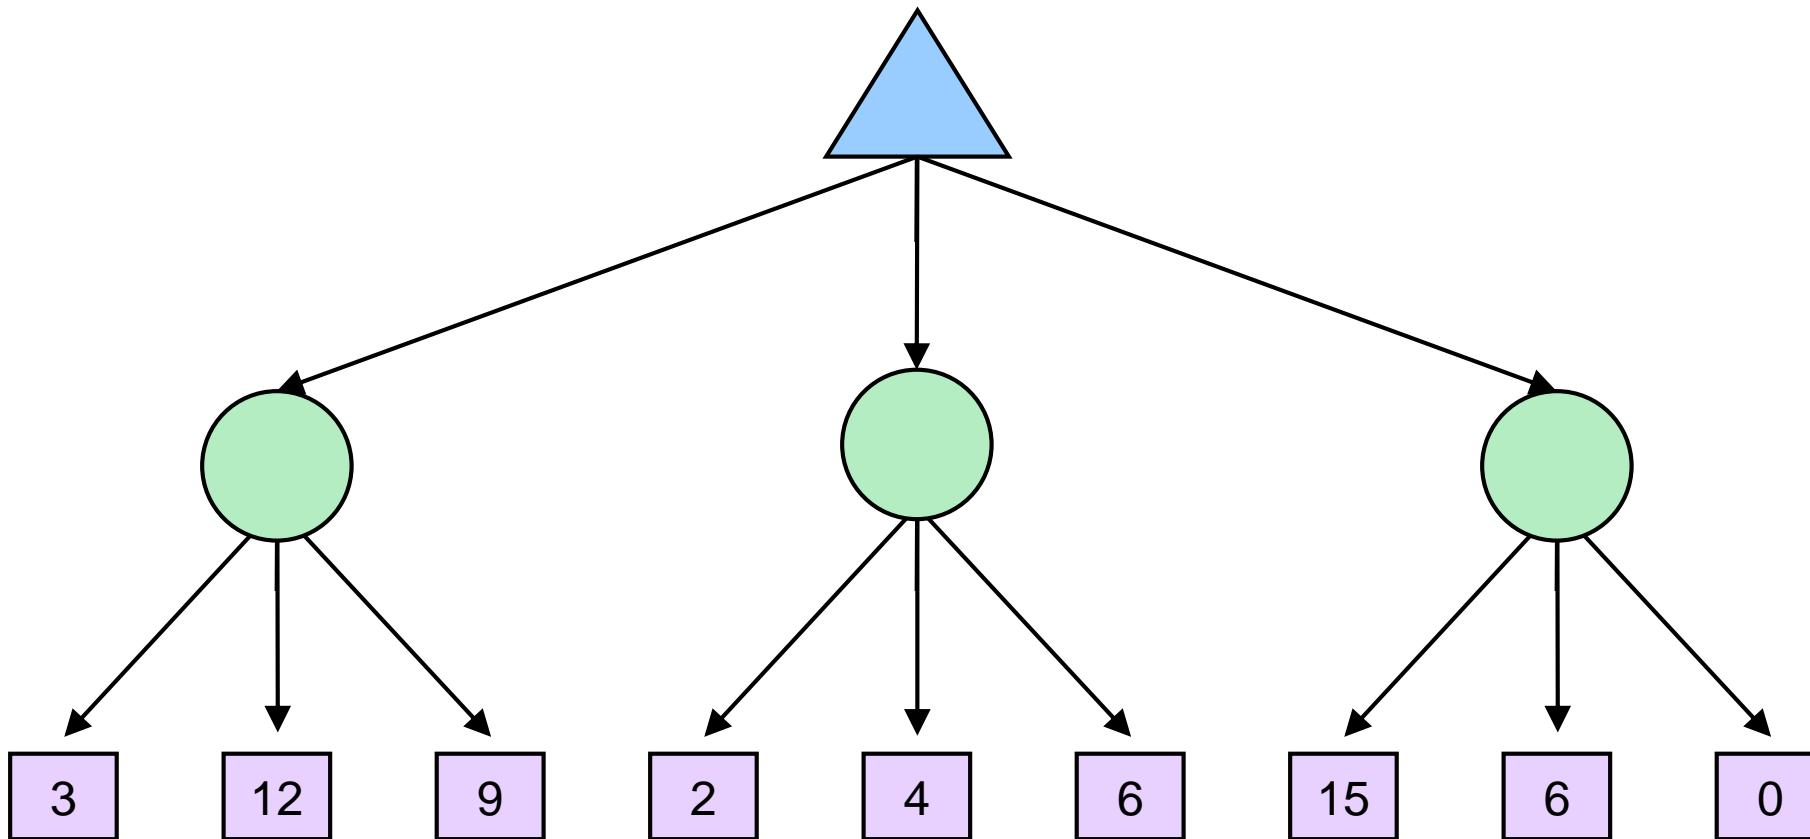
# Expectimax Pseudocode

```
def exp-value(state):  
    initialize v = 0  
    for each successor of state:  
        p = probability(successor)  
        v += p * value(successor)  
    return v
```

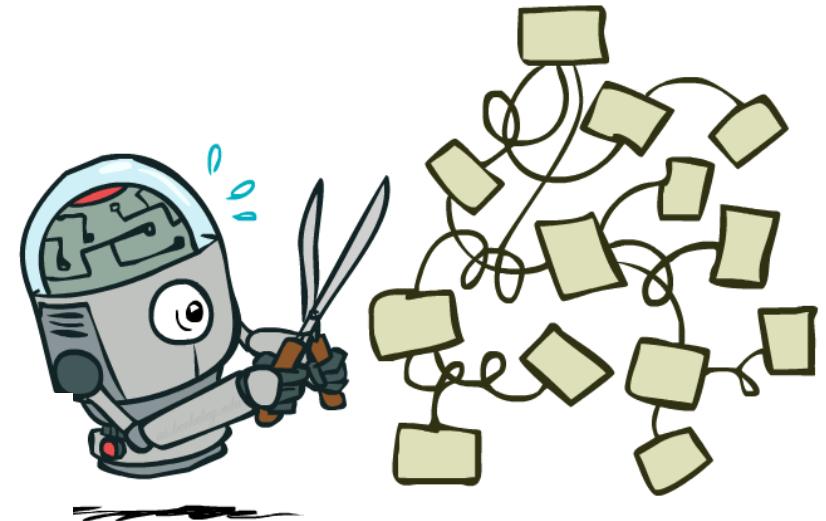
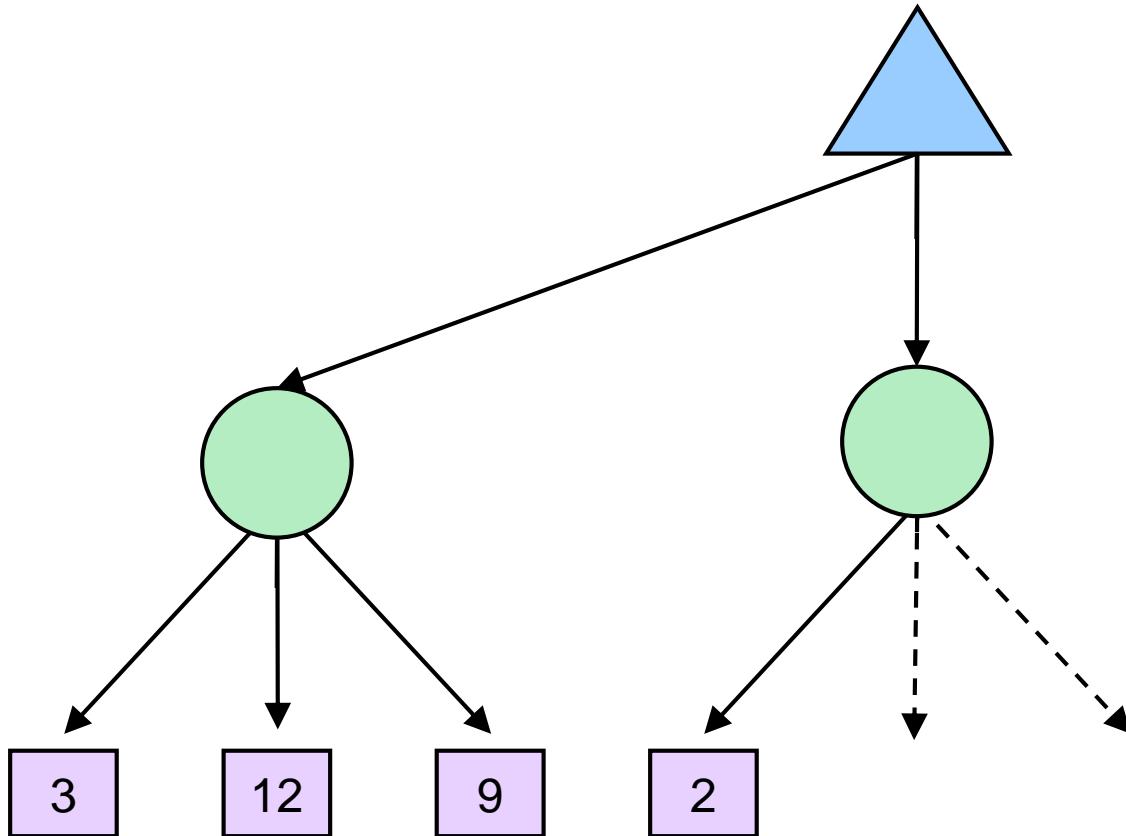


$$v = (1/2)(8) + (1/3)(24) + (1/6)(-12) = 10$$

# Expectimax Example

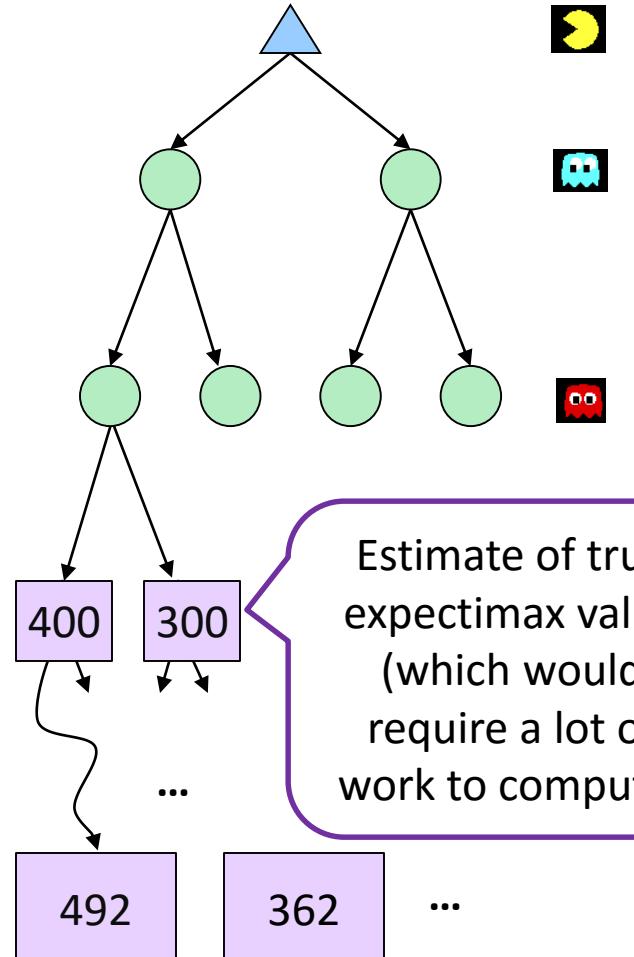


# Expectimax Pruning?



Can we Prune ?

# Depth-Limited Expectimax



can we depth Limited  
Search Here?