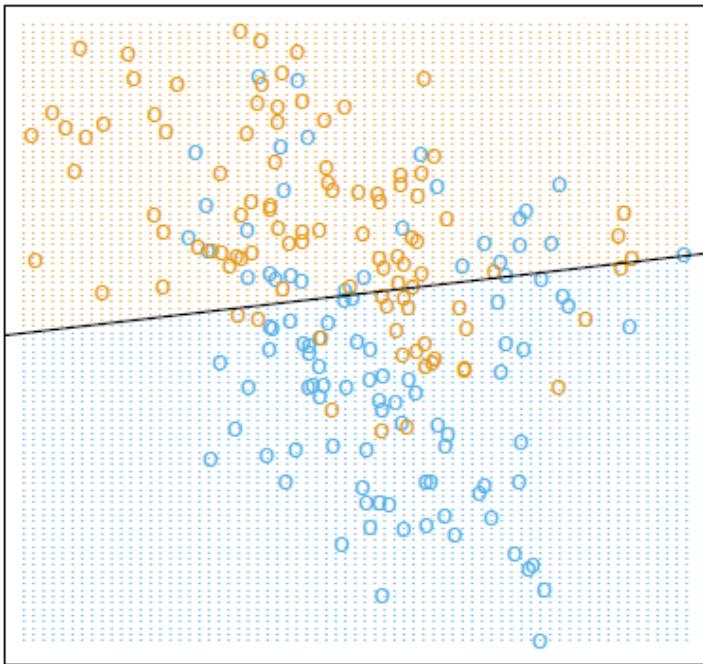


Linear Classification/
Logistic Regression/

Linear models for classification

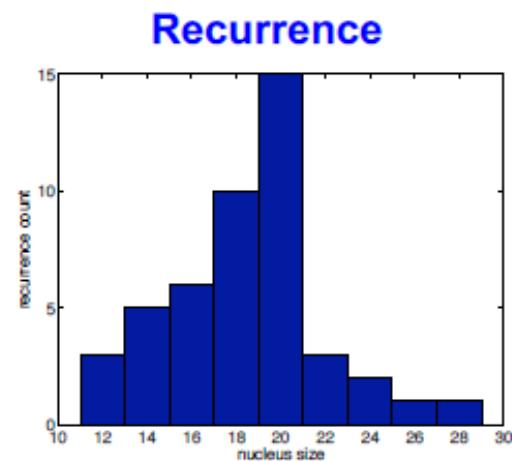
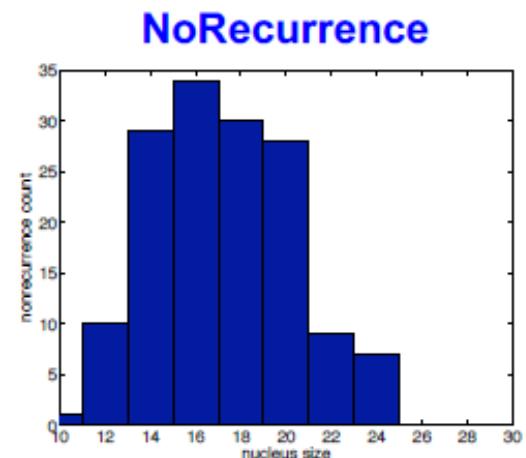
Linear Regression of 0/1 Response



- The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then fit by linear regression.
- The line is the decision boundary defined by $w^T x = 0.5$.
- The orange shaded region denotes that part of input space classified as ORANGE, while the blue region is classified as BLUE.

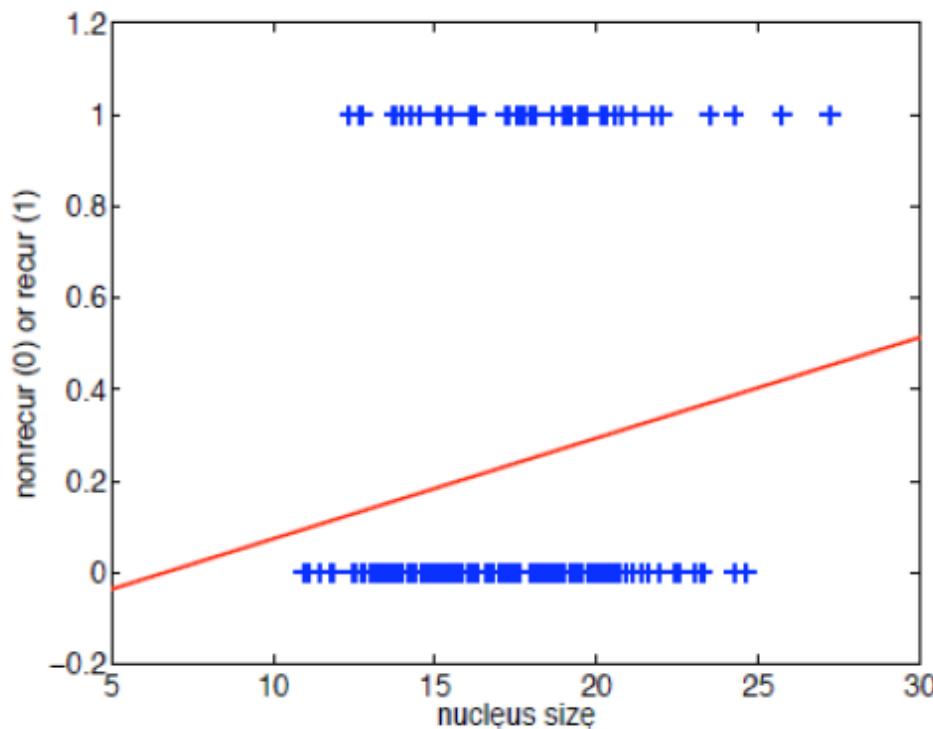
Simple example

- Given “nucleus size”, predict cancer recurrence.
- Univariate input: X = nucleus size.
- Binary output: $Y = \{\text{NoRecurrence} = 0; \text{Recurrence} = 1\}$
- Try: Minimize the least-square error.



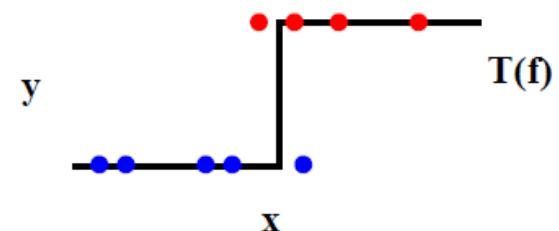
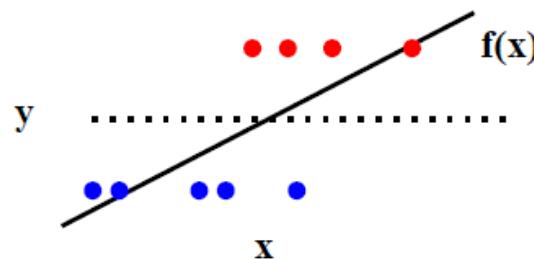
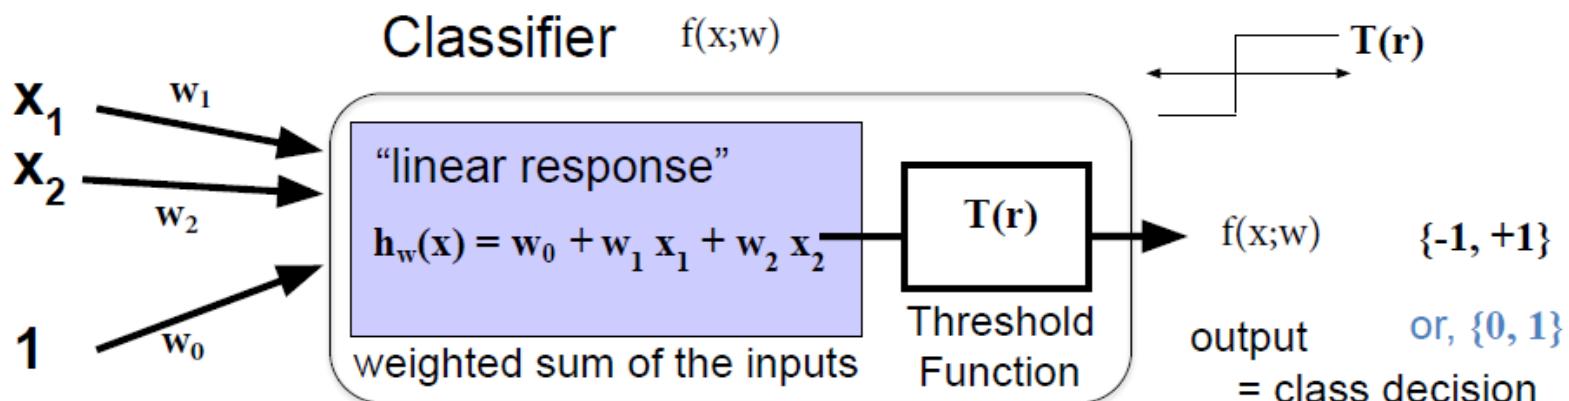
Predicting a class from linear regression

- How to get a binary output?
 1. Threshold the output:
 $\{ y \leq t \text{ for NoRecurrence, } y > t \text{ for Recurrence} \}$
 2. Interpret output as probability:
 $y = Pr(\text{Recurrence})$



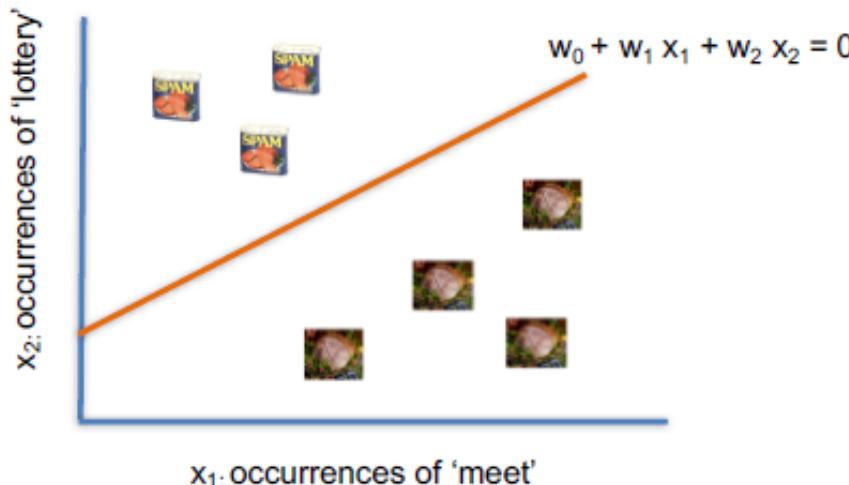
Perceptron

- Perceptron takes weighted linear combination of input features and pass it through a thresholding function which outputs 1 or 0.
- The sign of $w^T x$ tells us which side of the plane $w^T x = 0$, the point x lies on.
- Thus by taking threshold as 0, perceptron classifies data based on which side of the plane the new point lies on.



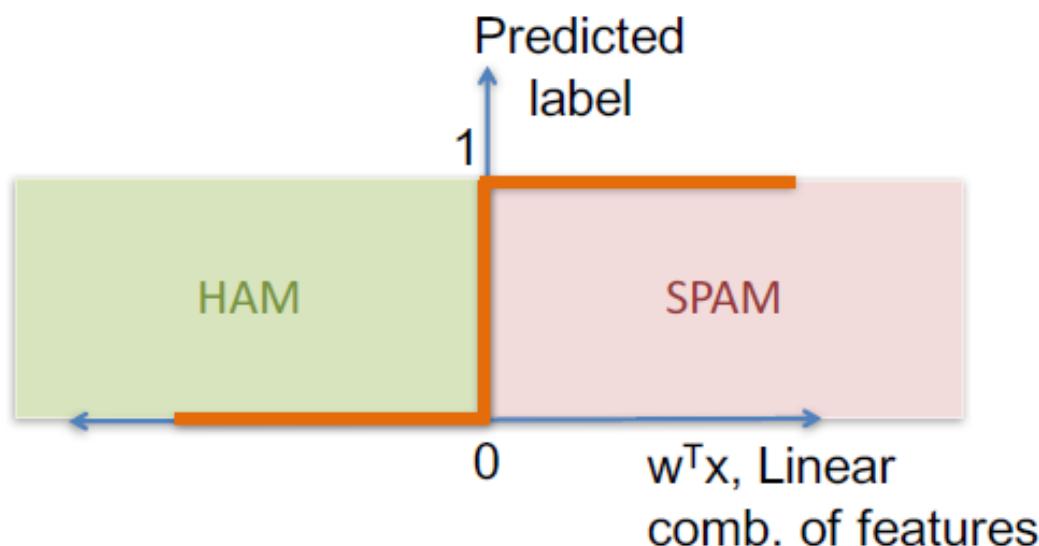
Example

- $x_1 = \#$ of times 'meet' appears in an email
- $x_2 = \#$ of times 'lottery' appears in an email
- Define feature vector $\mathbf{x} = [1, x_1, x_2]$
- Learn the decision boundary $w_0 + w_1 x_1 + w_2 x_2 = 0$ such that
 - If $\mathbf{w}^\top \mathbf{x} \geq 0$ declare $y = 1$ (spam)
 - If $\mathbf{w}^\top \mathbf{x} < 0$ declare $y = 0$ (ham)



Visualizing a Linear Classifier

- $x_1 = \#$ of times 'lottery' appears in an email
- $x_2 = \#$ of times 'meet' appears in an email
- Define feature vector $\mathbf{x} = [1, x_1, x_2]$
- Learn the decision boundary $w_0 + w_1x_1 + w_2x_2 = 0$ such that
 - If $\mathbf{w}^\top \mathbf{x} \geq 0$ declare $y = 1$ (spam)
 - If $\mathbf{w}^\top \mathbf{x} < 0$ declare $y = 0$ (ham)



$y = 1$ for spam, $y = 0$ for ham

Suppose you see the following email:

CONGRATULATIONS!! Your email address have won you the lottery sum of US\$2,500,000.00 USD to claim your prize, contact your office agent (Arthur walter) via email claims2155@yahoo.com.hk or call +44 704 575 1113

Keywords are [lottery, prize, office, email]

The given weight vector is $\mathbf{w} = [0.3, 0.3, -0.1, -0.04]^\top$

Will we predict that the email is spam or ham?

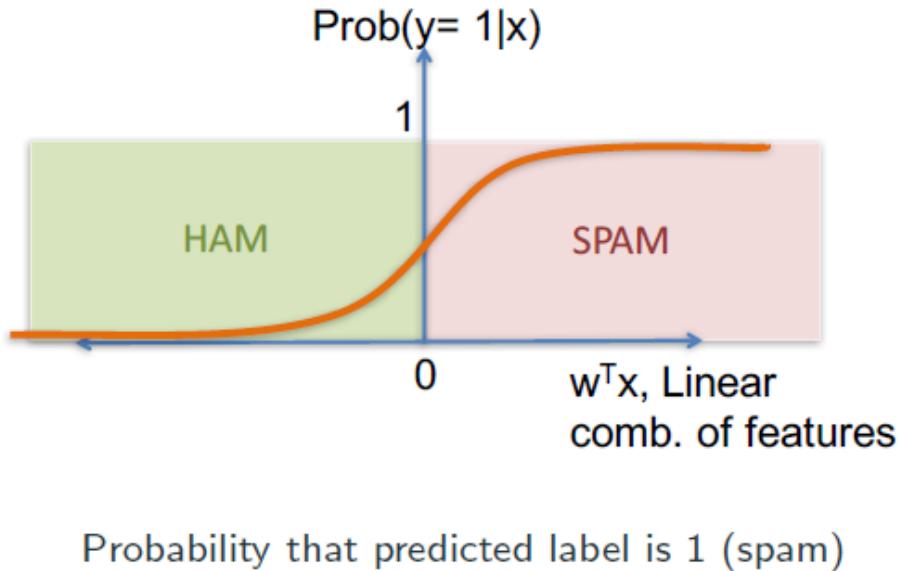
$$\mathbf{x} = [1, 1, 1, 2]^\top$$

$$\mathbf{w}^\top \mathbf{x} = 0.3 * 1 + 0.3 * 1 - 0.1 * 1 - 0.04 * 2 = 0.42 > 0$$

so we predict spam!

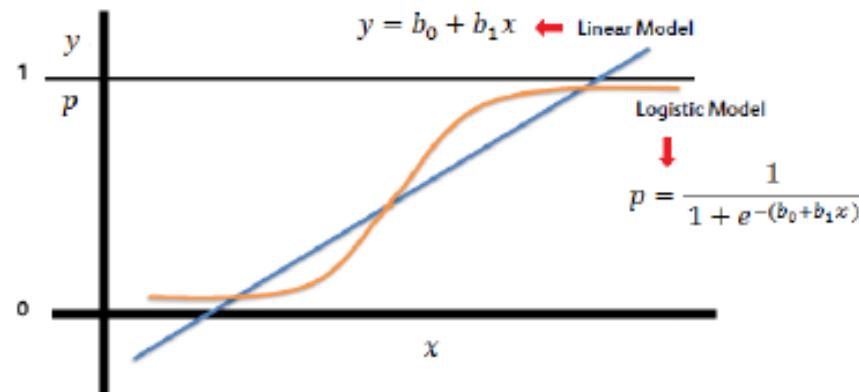
Intuition : Logistic Regression

- Suppose we want to output the **probability** of an email being spam/ham instead of just 0 or 1
- This gives information about the confidence in the decision
- Use a function $\sigma(\mathbf{w}^T \mathbf{x})$ that maps $\mathbf{w}^T \mathbf{x}$ to a value between 0 and 1

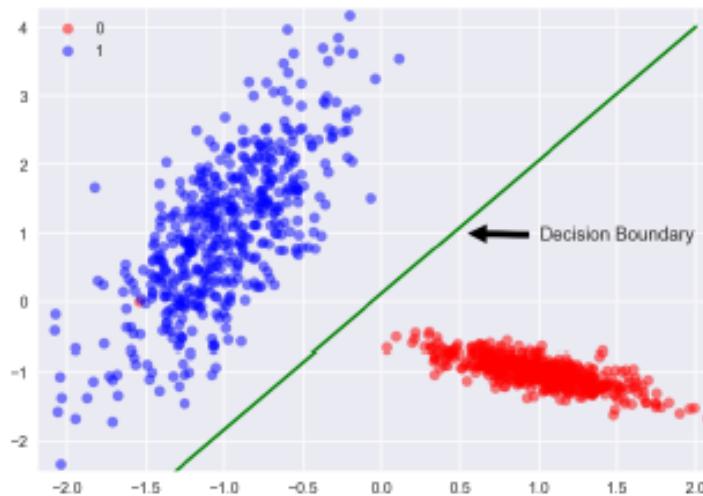


Key Problem: Finding optimal weights \mathbf{w} that accurately predict this probability for a new email

Sigmoid function returns values in [0,1]



Decision boundary is linear: **linear classifier**



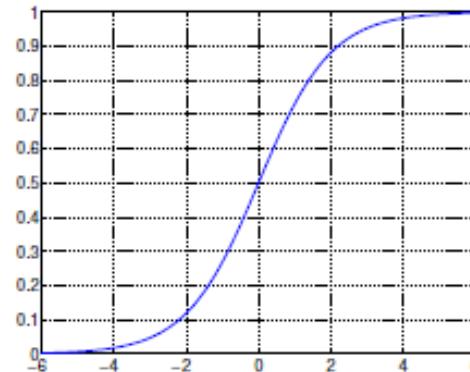
Why the Sigmoid Function?

What does it look like?

$$\sigma(a) = \frac{1}{1 + e^{-a}}$$

where

$$a = \mathbf{w}^\top \mathbf{x}$$



Sigmoid properties

- Bounded between 0 and 1 ← thus, interpretable as probability
- Monotonically increasing ← thus, usable to derive classification rules
 - $\sigma(a) \geq 0.5$, positive (classify as '1')
 - $\sigma(a) < 0.5$, negative (classify as '0')
- Nice computational properties ← as we will see soon

Suppose you see the following email:

CONGRATULATIONS!! Your email address have won you the lottery sum of US\$2,500,000.00 USD to claim your prize, contact your office agent (Arthur walter) via email claims2155@yahoo.com.hk or call +44 704 575 1113

Keywords are [lottery, prize, office, email]

The given weight vector is $\mathbf{w} = [0.3, 0.3, -0.1, -0.04]^\top$

What is the probability that the email is spam?

$$\mathbf{x} = [1, 1, 1, 2]^\top$$

$$\mathbf{w}^\top \mathbf{x} = 0.3 * 1 + 0.3 * 1 - 0.1 * 1 - 0.04 * 2 = 0.42 > 0$$

$$\Pr(y = 1 | \mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + e^{-0.42}} = 0.603$$

Logistic Regression Objective Function

- Can't just use squared loss as in linear regression:

$$J(\mathbf{w}) = \frac{1}{2n} \sum_{i=1}^n \left(h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)} \right)^2$$

- Using the logistic regression model

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

results in a non-convex optimization

Deriving the Cost Function via Maximum Likelihood Estimation

- The intuition behind it is to **find the parameters of a model that maximize the probability** of observing the given data.
- Product of the individual likelihoods for each data point when the data points are assumed to be independent.

$$l(\boldsymbol{\theta}) = \prod_{i=1}^n p(y^{(i)} \mid \mathbf{x}^{(i)}; \boldsymbol{\theta})$$

- Likelihood of data is given by:
- So, looking for the $\boldsymbol{\theta}$ that maximizes the likelihood

$$\boldsymbol{\theta}_{\text{MLE}} = \arg \max_{\boldsymbol{\theta}} l(\boldsymbol{\theta}) = \arg \max_{\boldsymbol{\theta}} \prod_{i=1}^n p(y^{(i)} \mid \mathbf{x}^{(i)}; \boldsymbol{\theta})$$

Deriving the Cost Function via Maximum Likelihood Estimation

- Can take the log without changing the solution:

$$\boldsymbol{\theta}_{\text{MLE}} = \arg \max_{\boldsymbol{\theta}} \log \prod_{i=1}^n p(y^{(i)} \mid \mathbf{x}^{(i)}; \boldsymbol{\theta})$$

$$= \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^n \log p(y^{(i)} \mid \mathbf{x}^{(i)}; \boldsymbol{\theta})$$

Deriving the Cost Function via Maximum Likelihood

Let us assume that

$$\begin{aligned} P(y = 1 \mid x; \theta) &= h_{\theta}(x) \\ P(y = 0 \mid x; \theta) &= 1 - h_{\theta}(x) \end{aligned}$$

Note that this can be written more compactly as

$$p(y \mid x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$$

Assuming that the m training examples were generated independently, we can then write down the likelihood of the parameters as

$$\begin{aligned} L(\theta) &= p(\vec{y} \mid X; \theta) \\ &= \prod_{i=1}^m p(y^{(i)} \mid x^{(i)}; \theta) \\ &= \prod_{i=1}^m (h_{\theta}(x^{(i)}))^{y^{(i)}} (1 - h_{\theta}(x^{(i)}))^{1-y^{(i)}} \end{aligned}$$

Deriving the Cost Function via Maximum Likelihood

- Expand as follows:

$$\theta_{\text{MLE}} = \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^n \log p(y^{(i)} | \mathbf{x}^{(i)}; \boldsymbol{\theta})$$

$$= \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^n \left[y^{(i)} \log p(y^{(i)} = 1 | \mathbf{x}^{(i)}; \boldsymbol{\theta}) + (1 - y^{(i)}) \log (1 - p(y^{(i)} = 1 | \mathbf{x}^{(i)}; \boldsymbol{\theta})) \right]$$

- Substitute in model, and take negative to yield

Logistic regression objective:

$$\min_{\boldsymbol{\theta}} J(\boldsymbol{\theta})$$

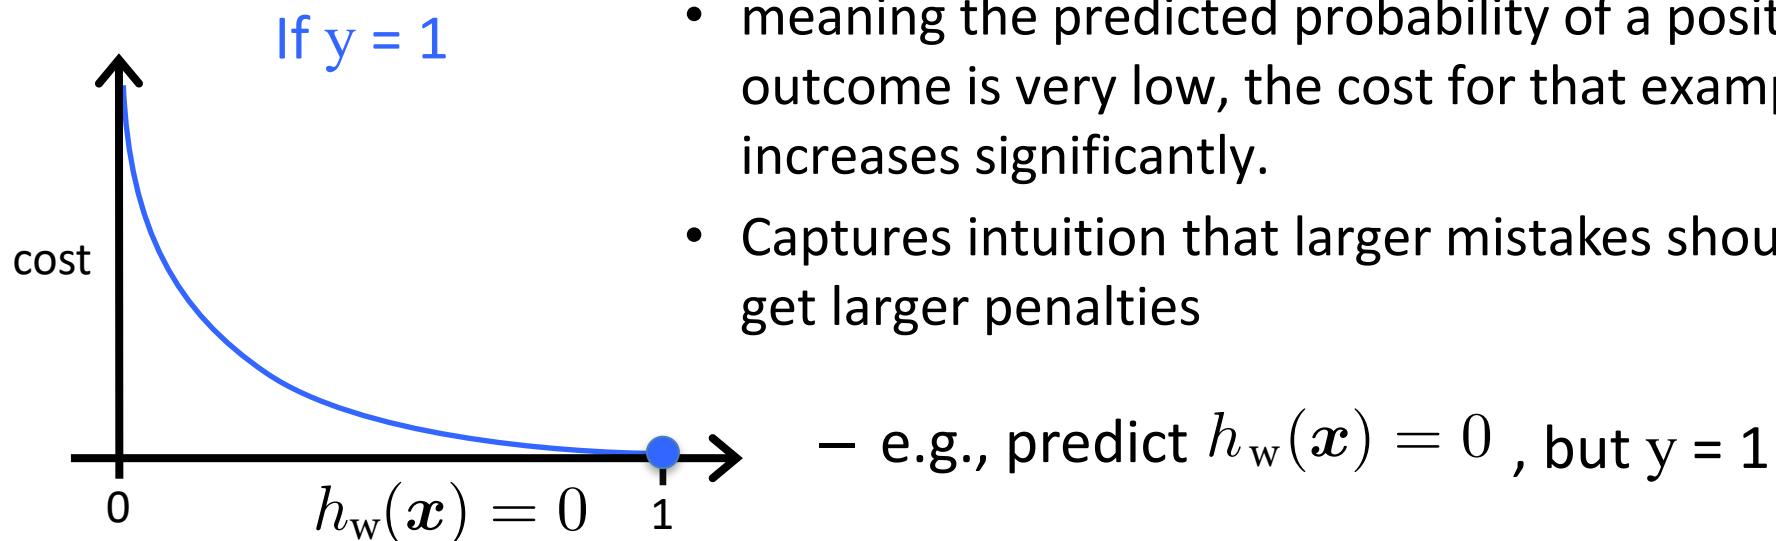
$$J(\boldsymbol{\theta}) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)})) \right]$$

Intuition Behind the Objective

$$\text{cost}(h_w(\mathbf{x}), y) = \begin{cases} -\log(h_w(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_w(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

When actual label y is 1

- Cost = 0 if prediction is correct
- As $h_w(\mathbf{x}) \rightarrow 0$, cost $\rightarrow \infty$
- meaning the predicted probability of a positive outcome is very low, the cost for that example increases significantly.
- Captures intuition that larger mistakes should get larger penalties



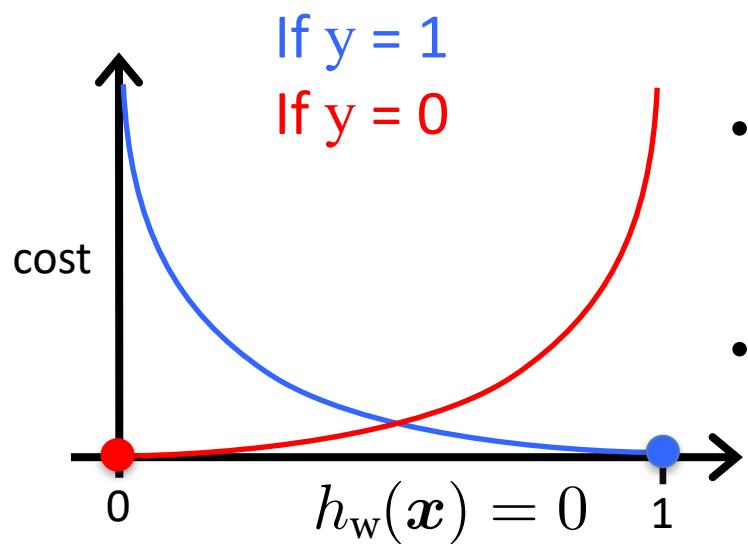
Intuition Behind the Objective

$$\text{cost}(h_w(\mathbf{x}), y) = \begin{cases} -\log(h_w(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_w(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

If $y = 0$

- Cost = 0 if prediction is correct
- As $(1 - h_w(\mathbf{x})) \rightarrow 0$, cost $\rightarrow \infty$

- meaning the predicted probability of a negative outcome is very low), the cost for that example increases significantly.
- Captures intuition that larger mistakes should get larger penalties



Gradient Descent for Logistic Regression

$$J_{\text{reg}}(\mathbf{w}) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\mathbf{w}}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\mathbf{w}}(\mathbf{x}^{(i)})) \right]$$

Want $\min_{\mathbf{w}} J(\mathbf{w})$

- Initialize \mathbf{w}
- Repeat until convergence

$$\mathbf{w}_j = \mathbf{w}_j - \alpha \frac{\partial}{\partial w_j} J(\mathbf{w})$$

simultaneous update
for $j = 0 \dots d$

Use the natural logarithm ($\ln = \log_e$) to cancel with the $\exp()$ in $h_{\mathbf{w}}(\mathbf{x})$

Gradient Descent for Logistic Regression

$$J_{\text{reg}}(\mathbf{w}) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\mathbf{w}}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\mathbf{w}}(\mathbf{x}^{(i)})) \right]$$

Want $\min_{\mathbf{w}} J(\mathbf{w})$

- Initialize \mathbf{w}
- Repeat until convergence (simultaneous update for $j = 0 \dots d$)

$$\mathbf{w}_0 = \mathbf{w}_0 - \alpha \sum_{i=1}^n \left(h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)} \right)$$

$$\mathbf{w}_j = \mathbf{w}_j - \alpha \left[\sum_{i=1}^n \left(h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)} \right) x_j^{(i)} \right]$$

Gradient Descent for Logistic Regression

- Initialize w
- Repeat until convergence (simultaneous update for $j = 0 \dots d$)

$$w_0 \leftarrow w_0 - \alpha \sum_{i=1}^n \left(h_w(\mathbf{x}^{(i)}) - y^{(i)} \right)$$

$$w_j \leftarrow w_j - \alpha \left[\sum_{i=1}^n \left(h_w(\mathbf{x}^{(i)}) - y^{(i)} \right) x_j^{(i)} \right]$$

This looks IDENTICAL to linear regression!!!

- Ignoring the $1/n$ constant
- However, the form of the model is very different:

$$h_w(\mathbf{x}) = \frac{1}{1 + e^{-w^\top \mathbf{x}}}$$

Gradient Descent for Logistic Regression

$$J_{\text{reg}}(\boldsymbol{\theta}) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)})) \right] + \lambda \|\boldsymbol{\theta}_{[1:d]}\|_2^2$$

Want $\min_{\boldsymbol{\theta}} J(\boldsymbol{\theta})$

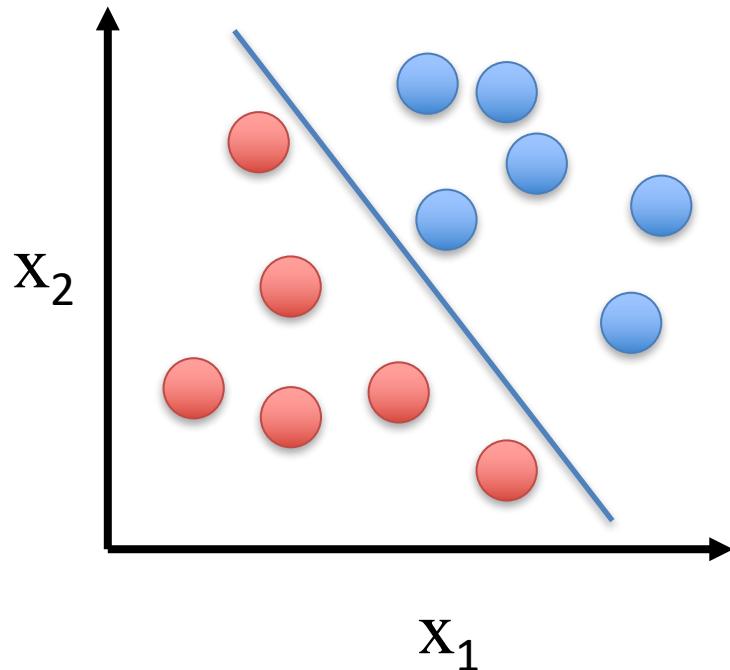
- Initialize $\boldsymbol{\theta}$
- Repeat until convergence (simultaneous update for $j = 0 \dots d$)

$$\theta_0 \quad \theta_0 - \alpha \sum_{i=1}^n \left(h_{\boldsymbol{\theta}} \left(\mathbf{x}^{(i)} \right) - y^{(i)} \right)$$

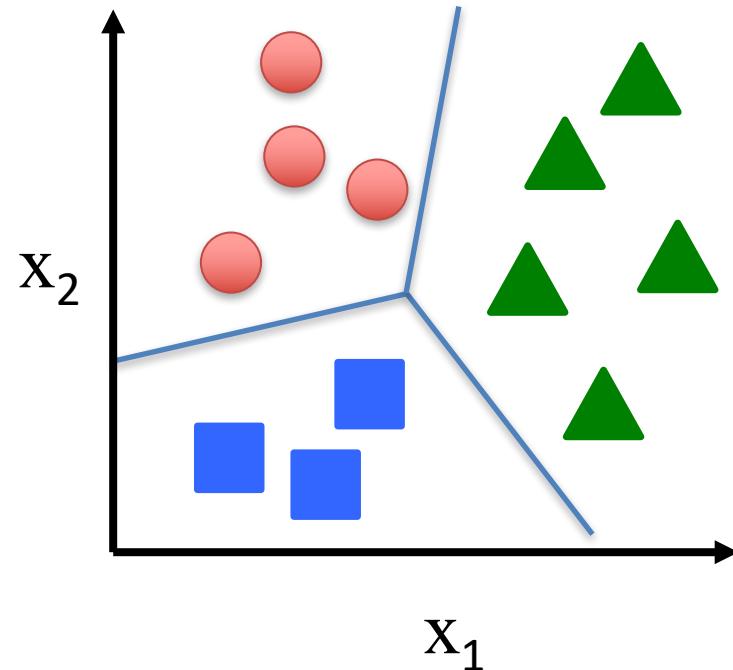
$$\theta_j \quad \theta_j - \alpha \left[\sum_{i=1}^n \left(h_{\boldsymbol{\theta}} \left(\mathbf{x}^{(i)} \right) - y^{(i)} \right) x_j^{(i)} - \frac{\lambda}{n} \theta_j \right]$$

Multi-Class Classification

Binary classification:



Multi-class classification:

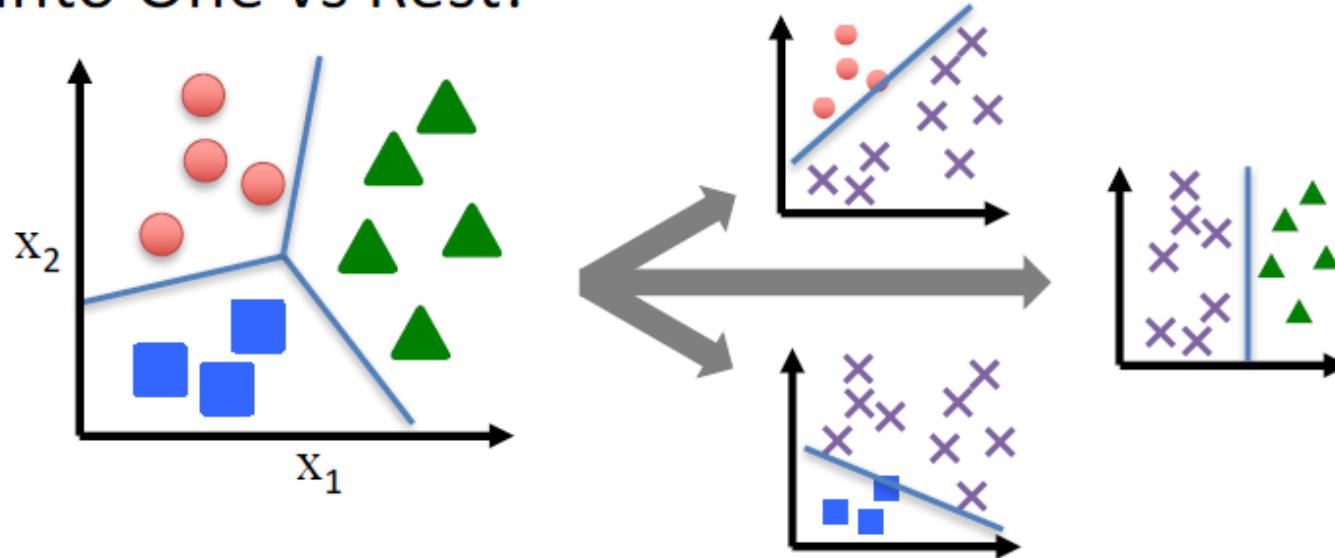


Disease diagnosis: healthy / cold / flu / pneumonia

Object classification: desk / chair / monitor / bookcase

Multi-Class Logistic Regression

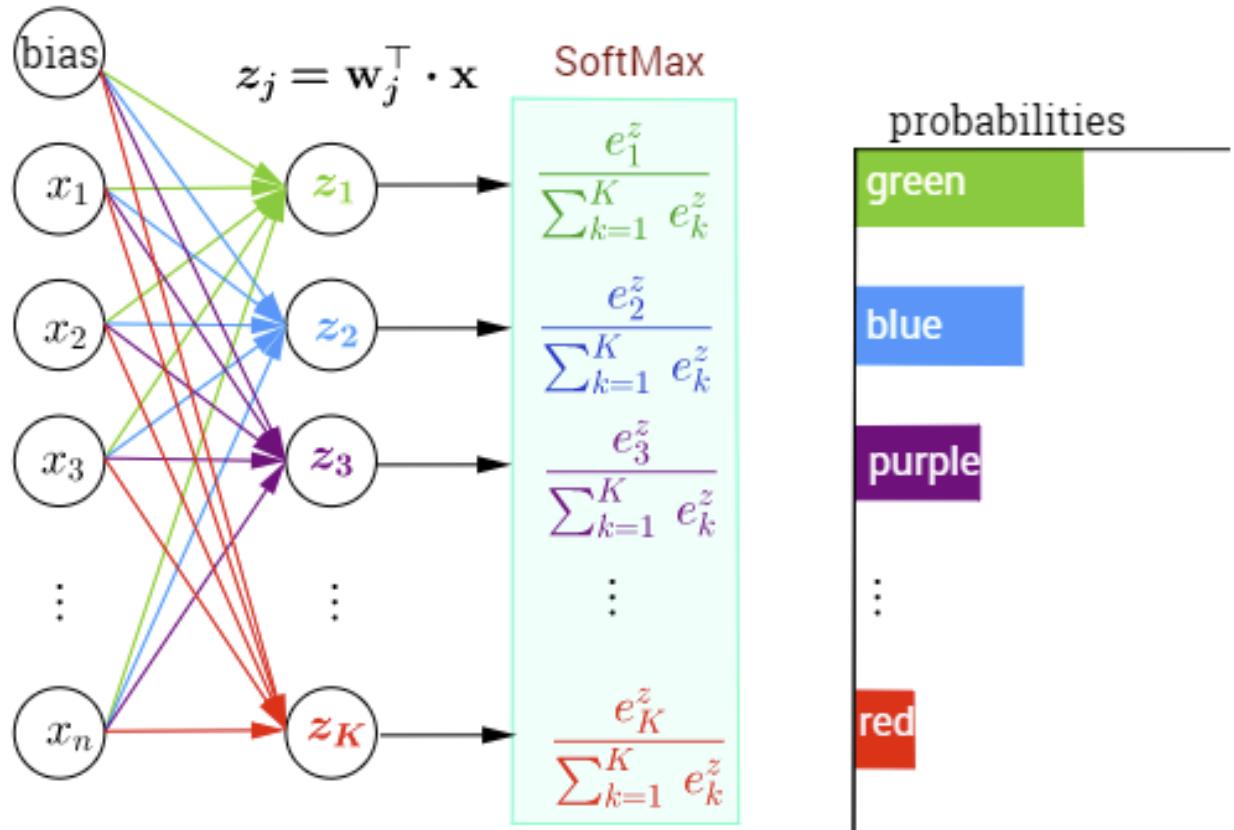
Split into One vs Rest:



- Train a logistic regression classifier for each class i to predict the probability that $y = i$ with

$$h_c(\mathbf{x}) = \frac{\exp(\boldsymbol{\theta}_c^\top \mathbf{x})}{\sum_{c=1}^C \exp(\boldsymbol{\theta}_c^\top \mathbf{x})}$$

$$\mathbf{z} = \begin{bmatrix} z_1 \\ z_2 \\ z_3 \\ \vdots \\ z_K \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^\top \\ \mathbf{w}_2^\top \\ \mathbf{w}_3^\top \\ \vdots \\ \mathbf{w}_K^\top \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}$$



Multi-Class Logistic Regression

- For 2 classes:

$$h_{\theta}(x) = \frac{1}{1 + \exp(-\theta^T x)} = \frac{\exp(\theta^T x)}{1 + \exp(\theta^T x)}$$

- For C classes {1, ..., C}:

$$p(y = c \mid x; \theta_1, \dots, \theta_C) = \frac{\exp(\theta_c^T x)}{\sum_{c=1}^C \exp(\theta_c^T x)}$$

– Called the **softmax** function

Implementing Multi-Class Logistic Regression

- Use $h_c(\mathbf{x}) = \frac{\exp(\boldsymbol{\theta}_c^\top \mathbf{x})}{\sum_{c=1}^C \exp(\boldsymbol{\theta}_c^\top \mathbf{x})}$ as the model for class c
- Gradient descent simultaneously updates all parameters for all models
 - Same derivative as before, just with the above $h_c(\mathbf{x})$
- Predict class label as the most probable label

$$\max_c h_c(\mathbf{x})$$